

[Sign In](#)[DSA](#) [Practice Problems](#) [C](#) [C++](#) [Java](#) [Python](#) [JavaScript](#) [Data Science](#) [Machine Learning](#) [Courses](#)

# RMSProp Optimizer in Deep Learning



Last Updated : 30 Sep, 2025

RMSProp (Root Mean Square Propagation) is an adaptive learning rate optimization algorithm designed to improve the performance and speed of training deep learning models.

- It is a variant of the [gradient descent](#) algorithm which adapts the learning rate for each parameter individually by considering the magnitude of recent gradients for those parameters.
- This adaptive nature helps in dealing with the challenges of non-stationary objectives and sparse gradients commonly encountered in deep learning tasks.


**Escuela de Psicología  
CETYS Universidad  
Pacific Oaks College**

The three-day, international conference explored views on trauma and culture changes individual and social.

## Need of RMSProp Optimizer

RMSProp was developed to address the limitations of previous optimization methods such as [SGD](#)

[Skip to content](#)



## (Stochastic Gradient Descent) and AdaGrad as SGD

uses a constant learning rate which can be inefficient and AdaGrad reduces the learning rate too aggressively.

RMSProp balances by adapting the learning rates based on a moving average of squared gradients. This approach helps in maintaining a balance between efficient convergence and stability during the training process making RMSProp a widely used optimization algorithm in modern deep learning.

## How RMSProp Works?

RMSProp keeps a moving average of the squared gradients to normalize the gradient updates. By doing so it prevents the learning rate from becoming too small which was a drawback in AdaGrad and ensures that the updates are appropriately scaled for each parameter. This mechanism allows RMSProp to perform well even in the presence of non-stationary objectives making it suitable for training deep learning models.

The mathematical formulation is as follows:

1. Compute the gradient  $g_t$  at time step t:

[Skip to content](#)

$$gt = \nabla\theta$$

2. Update the moving average of squared gradients:

$$E[g^2]_t = \gamma E[g^2]_{t-1} + (1 - \gamma)$$

where  $\gamma$  is the decay rate.

3. Update the parameter  $\theta$  using the adjusted learning rate:

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}}$$

where  $\eta$  is the learning rate and  $\epsilon$  is a small constant added for numerical stability.

## Parameters Used in RMSProp

- **Learning Rate ( $\eta$ ):** Controls the step size during the parameter updates. RMSProp typically uses a default learning rate of 0.001, but it can be adjusted based on the specific problem.
- **Decay Rate ( $\gamma$ ):** Determines how quickly the moving average of squared gradients decays. A

Skip to content

common default value is 0.9 which balances the contribution of recent and past gradients.

- **Epsilon ( $\epsilon$ ):** A small constant added to the denominator to prevent division by zero and ensure numerical stability. A typical value for  $\epsilon$  is 1e-8.

By carefully adjusting these parameters, RMSProp effectively adapts the learning rates during training, leading to faster and more reliable convergence in deep learning models.

## Implementing RMSprop in Python using TensorFlow or Keras

We will use the following code line for initializing the RMSProp optimizer with hyperparameters:

```
tf.keras.optimizers.RMSprop(learning_rate=0.001, rho=0.9)
```

- **learning\_rate=0.001:** Sets the step size for weight updates. Smaller learning rates result in smaller updates, helping to fine-tune weights and prevent overshooting the minimum loss.

[Skip to content](#)

- **rho=0.9:** The discounting factor for the history of gradients, controlling the influence of past gradients on the current gradient computation.

## 1. Importing Libraries

We are importing libraries to implement RMSprop optimizer, handle datasets, build the model and plot results.

- tensorflow.keras for deep learning components.
- matplotlib.pyplot for visualization.

```
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense,
Flatten
from tensorflow.keras.utils import
to_categorical
import matplotlib.pyplot as plt
```

## 2. Loading and Preprocessing Dataset

We load the MNIST dataset, normalize pixel values to [0,1] and [one-hot encode labels](#).

- **mnist.load\_data()** loads images and labels.

Skip to content

- **Normalization** improves training stability.
- **to\_categorical()** converts labels to one-hot vectors.

```
(x_train, y_train), (x_test, y_test) =   
mnist.load_data()  
  
x_train = x_train.astype('float32') / 255.0  
x_test = x_test.astype('float32') / 255.0  
y_train = to_categorical(y_train, 10)  
y_test = to_categorical(y_test, 10)
```

### 3. Building the Model

We define a neural network using Sequential with input flattening and dense layers.

- Flatten converts 2D images to 1D vectors.
- Dense layers learn patterns with ReLU and softmax activations.

```
model = Sequential([  
    Flatten(input_shape=(28, 28)),  
    Dense(128, activation='relu'),  
    Dense(64, activation='relu'),  
    Dense(10, activation='softmax')  
]) 
```

[Skip to content](#)

## 4. Compiling the Model

We compile the model using the RMSprop optimizer for adaptive learning rates, categorical cross-entropy loss for multi-class classification and track accuracy metric.

- RMSprop adjusts learning rates based on recent gradients (parameter rho controls decay rate).
- categorical\_crossentropy suits one-hot encoded labels.

```
model.compile(optimizer=tf.keras.optimizers.RM ↴  
rop(learning_rate=0.001, rho=0.9),  
          loss='categorical_crossentropy',  
          metrics=['accuracy'])
```

## 5. Training the Model

We train the model over 10 epochs with batch size 32 and validate on 20% of training data. validation\_split monitors model performance on unseen data each epoch.

```
history = model.fit(x_train, y_train, epochs=1 ↴  
                     batch_size=32,  
                     validation_split=0.2)
```

Skip to content

**Output:**

```

Epoch 1/10
1500/1500    7s 4ms/step - accuracy: 0.8721 - loss: 0.4363 - val_accuracy: 0.9524 - val_loss: 0.1541
Epoch 2/10
1500/1500    9s 4ms/step - accuracy: 0.9624 - loss: 0.1212 - val_accuracy: 0.9673 - val_loss: 0.1173
Epoch 3/10
1500/1500    6s 4ms/step - accuracy: 0.9737 - loss: 0.0864 - val_accuracy: 0.9685 - val_loss: 0.1177
Epoch 4/10
1500/1500    6s 4ms/step - accuracy: 0.9799 - loss: 0.0664 - val_accuracy: 0.9719 - val_loss: 0.1088
Epoch 5/10
1500/1500    6s 4ms/step - accuracy: 0.9838 - loss: 0.0526 - val_accuracy: 0.9747 - val_loss: 0.0992
Epoch 6/10
1500/1500    10s 4ms/step - accuracy: 0.9880 - loss: 0.0397 - val_accuracy: 0.9737 - val_loss: 0.1076
Epoch 7/10
1500/1500    7s 5ms/step - accuracy: 0.9907 - loss: 0.0338 - val_accuracy: 0.9736 - val_loss: 0.1200
Epoch 8/10
1500/1500    10s 5ms/step - accuracy: 0.9920 - loss: 0.0284 - val_accuracy: 0.9749 - val_loss: 0.1208
Epoch 9/10
1500/1500    6s 4ms/step - accuracy: 0.9928 - loss: 0.0242 - val_accuracy: 0.9787 - val_loss: 0.1409
Epoch 10/10
1500/1500    11s 4ms/step - accuracy: 0.9933 - loss: 0.0214 - val_accuracy: 0.9758 - val_loss: 0.1398

```

*Training the Model*

## 6. Evaluating and Visualizing Results

We evaluate test accuracy on unseen test data and plot training and validation loss curves to visualize learning progress.


```

loss, accuracy = model.evaluate(x_test, y_test)
print(f'Test accuracy: {accuracy:.4f}')

plt.plot(history.history['loss'],
         label='Training Loss')
plt.plot(history.history['val_loss'],
         label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Cost Function Graph')
plt.legend()
plt.show()

```

**Output:***Skip to content*



Evaluating and Visualizing Results

## Advantages

- **Adaptive Learning Rates:** Adjusts learning rates for each parameter individually, optimizing updates more effectively.
- **Handles Non-Stationary Objectives:** Efficiently adapts to changing optimal parameter values over time.
- **Prevents Learning Rate Decay Problem:** Maintains optimal learning rates by using a decay rate unlike AdaGrad.

Skip to content

- **Improved Convergence Speed:** Faster convergence due to balanced and dynamic learning rates.

## Disadvantages

- **Sensitivity to Hyperparameters:** Performance is sensitive to settings like decay rate and epsilon meaning it requires careful tuning.
- **Poor Performance with Sparse Data:** May struggle with sparse data, leading to slower or inconsistent convergence.

### Suggested Quiz

5 Questions

Which of the following best describes the purpose of RMSProp in optimization?

- A It uses a fixed learning rate for all parameters.
- B It ignores past gradients and only uses the current gradient for updates.
- C It adapts the learning rate for each parameter based on recent gradients

[Skip to content](#)

[Login to View Explanation](#)

1/5

&lt; Previous

Next &gt;

[Comment](#) alka1974 + Follow

3

### Article Tags :

[Blogathon](#)    [Deep Learning](#)    [AI-ML-DS](#)[AI-ML-DS With Python](#)[+1 More](#)

## Explore

[Deep Learning Basics](#)[Neural Networks Basics](#)[Deep Learning Models](#)[Deep Learning Frameworks](#)[Model Evaluation](#)[Skip to content](#)

## Deep Learning Projects

---



**Corporate & Communications Address:**  
A-143, 7th Floor, Sovereign Corporate Tower, Sector- 136, Noida, Uttar Pradesh (201305)

**Registered Address:**  
K 061, Tower K, Gulshan Vivante Apartment, Sector 137, Noida, Gautam Buddh Nagar, Uttar Pradesh, 201305

Company	Explore	Tutorials	Courses	Videos	Preparation Corner
About Us	POTD	Programming	ML and Data	DSA	Interview Corner
Legal	Job-A-Thon	Languages	Science	Python	Aptitude
Privacy Policy	Blogs	DSA	DSA and	Java	Puzzles
Contact Us	Nation Skill Up	Web Technology	Placements	C++	GfG 160
Advertise with us		AI, ML & Data	Web	Web	System Design
GFG Corporate		Science	Development	Development	
Solution		DevOps	Programming	Data Science	
Campus Training		CS Core Subjects	Languages	CS Subjects	
Program		Interview	DevOps & Cloud		
		Preparation	GATE		
		Software and	Trending		
		Tools	Technologies		



@GeeksforGeeks, Sanchhaya Education Private Limited, All rights reserved

Do Not Sell or Share My Personal Information